

Language Technology

Chapter 8: Part-of-Speech Tagging Using Machine-Learning Techniques

Pierre Nugues

`Pierre.Nugues@cs.lth.se`

September 15 and 22, 2022



Motivation

The analysis of sentences often involves the analysis of words.

We can divide it in three main tasks:

- 1 Identify the type of word, for instance noun or verb using the classical grammar;
- 2 Identify a group or segment, a noun group for instance:
The big table,
- 3 Identify a name (a proper noun) for instance are these three words,
Kjell Olof Andersson,

This lecture will show you how to solve the first one, part-of-speech tagging, and you will write a program for the second one, chunking, in a next laboratory assignment.



Part-of-Speech Annotation (CoNLL 2000)

Annotation of: *He reckons the current account deficit will narrow to only # 1.8 billion in September.* We set aside the last column for now.

He	PRP	B-NP
reckons	VBZ	B-VP
the	DT	B-NP
current	JJ	I-NP
account	NN	I-NP
deficit	NN	I-NP
will	MD	B-VP
narrow	VB	I-VP
to	TO	B-PP
only	RB	B-NP
#	#	I-NP
1.8	CD	I-NP
billion	CD	I-NP
in	IN	B-PP
September	NNP	B-NP
.	.	O



Designing a Part-of-Speech Tagger

We will now create part-of-speech taggers, where we will examine three architectures:

- ① A feed-forward pipeline with a one-hot encoding of the words;
- ② A feed-forward pipeline with word embeddings: We will replace the one-hot vectors with GloVe embeddings;
- ③ A recurrent neural network, either a simple RNN or a LSTM, with word embeddings.



Features for Part-of-Speech Tagging

The word *visit* is ambiguous in English:

*I paid a **visit** to a friend* → noun

*I went to **visit** a friend* → verb

The context of the word enables us to tell, here an article or the infinitive marker

To train and apply the model, the tagger extracts a set of features from the surrounding words, for example, a sliding window spanning five words and centered on the current word.

We then associate the feature vector $(w_{i-2}, w_{i-1}, w_i, w_{i+1}, w_{i+2})$ with the part-of-speech tag t_i at index i .



Training Set

Part-of-speech taggers use a training set where every word is hand-annotated (Penn Treebank and CoNLL 2008).

Index	Word	Hand annotation	Index	Word	Hand annotation
1	Battle	JJ	19	of	IN
2	-	HYPH	20	their	PRP\$
3	tested	JJ	21	countrymen	NNS
4	Japanese	JJ	22	to	TO
5	industrial	JJ	23	visit	VB
6	managers	NNS	24	Mexico	NNP
7	here	RB	25	,	,
8	always	RB	26	a	DT
9	buck	VBP	27	boatload	NN
10	up	RP	28	of	IN
11	nervous	JJ	29	samurai	FW
12	newcomers	NNS	30	warriors	NNS
13	with	IN	31	blown	VBN
14	the	DT	32	ashore	RB
15	tale	NN	33	375	CD
16	of	IN	34	years	NNS
17	the	DT	35	ago	RB
18	first	JJ	36	.	.



Architecture 1: Part-of-Speech Tagging with Linear Classifiers

Linear classifiers are efficient devices to carry out part-of-speech tagging:

- ❶ The lexical values are the input data to the tagger.
- ❷ The parts of speech are assigned from left to right by the tagger.

ID	FORM	PPOS	
	BOS	BOS	Padding
	BOS	BOS	
1	Battle	NN	
2	-	HYPH	
3	tested	NN	
...	
17	the	DT	
18	first	JJ	
19	of	IN	
20	their	PRP\$	
21	countrymen	NNS	Input features
22	to	TO	
23	visit	VB	Predicted tag
24	Mexico		
25	,		↓
26	a		
27	boatload		
...	
34	years		
35	ago		
36	.		
	EOS		Padding
	EOS		



Feed-Forward Structure

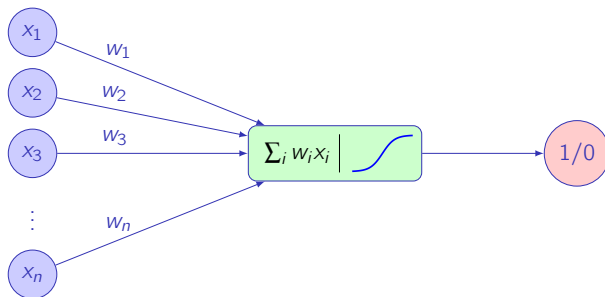
As input, the classifier uses:
 $(w_{i-2}, w_{i-1}, w_i, w_{i+1}, w_{i+2})$ to
 predict the part-of-speech tag t_i at
 index i .

Here:
 (countrymen, to, visit, Mexico, ",")
 to predict VB

ID	FORM	PPOS	
	BOS	BOS	Padding
	BOS	BOS	
1	Battle	NN	
2	-	HYPH	
3	tested	NN	
...	
17	the	DT	
18	first	JJ	
19	of	IN	
20	their	PRP\$	
21	countrymen	NNS	Input features
22	to	TO	
23	visit	VB	Predicted tag
24	Mexico		
25	,		↓
26	a		
27	boatload		
...	
34	years		
35	ago		
36	.		
	EOS		Padding
	EOS		




Feed Forward (Binary)



Logistic function



Feed Forward (Multilayer)




For the first layer, we have:

$$\text{activation}(\mathbf{W}^{(1)}\mathbf{x} + \mathbf{b}^{(1)}).$$



Feed Forward (Multinomial)



Softmax



Feature Vectors

ID	Feature vectors: X					PPOS: y
	w_{i-2}	w_{i-1}	w_i	w_{i+1}	w_{i+2}	
1	BOS	BOS	Battle	-	tested	NN
2	BOS	Battle	-	tested	Japanese	HYPH
3	Battle	-	tested	Japanese	industrial	JJ
...
19	the	first	of	their	countrymen	IN
20	first	of	their	countrymen	to	PRP\$
21	of	their	countrymen	to	visit	NNS
22	their	countrymen	to	visit	Mexico	TO
23	countrymen	to	visit	Mexico	,	VB
24	to	visit	Mexico	,	a	NNP
25	visit	Mexico	,	a	boatload	,
...
34	ashore	375	years	ago	.	NNS
35	375	years	ago	.	EOS	RB
36	years	ago	.	EOS	EOS	.




Word Encoding: One hot encoding

The feature space is defined by all the word values and a word has one dimension

This is a sparse representation

We use DictVectorizer() to encode them:

```
from sklearn.feature_extraction import DictVectorizer
v = DictVectorizer(sparse=False)
X_cat = [{ 'w_1': 'the', 'w_2': 'first', 'w_3': 'of' },
          { 'w_1': 'first', 'w_2': 'of', 'w_3': 'the' },
          { 'w_1': 'of', 'w_2': 'the', 'w_3': 'countrymen' },
          { 'w_1': 'the', 'w_2': 'countrymen', 'w_3': 'to' }]
X = v.fit_transform(X_cat)
X
array([[0., 0., 1., 0., 1., 0., 0., 0., 1., 0., 0.],
       [1., 0., 0., 0., 0., 1., 0., 0., 0., 0., 1.],
       [0., 1., 0., 0., 0., 0., 1., 1., 0., 0., 0.],
       [0., 0., 1., 1., 0., 0., 0., 0., 0., 0., 1.]])
```



Word Encoding

```
v.transform([{'w_1': 'the', 'w_2': 'of', 'w_3': 'Mexico'}])  
array([[0., 0., 1., 0., 0., 1., 0., 0., 0., 0.]])
```

```
v.get_feature_names()  
['w_1=first',  
 'w_1=of',  
 'w_1=the',  
 'w_2=countrymen',  
 'w_2=first',  
 'w_2=of',  
 'w_2=the',  
 'w_3=countrymen',  
 'w_3=of',  
 'w_3=the',  
 'w_3=to']
```

Notebook on GitHub:
[ilppp/programs/ch08/python/lr_pos_tagger_simple.ipynb](#)



Architecture 1: A Feed-Forward Neural Network

We use a Keras feed-forward architecture corresponding to a logistic regression:

```
np.random.seed(0)

model = models.Sequential([Dense(NB_CLASSES,
                                  input_dim=X.shape[1],
                                  activation='softmax')])

model.compile(loss='sparse_categorical_crossentropy',
              optimizer=OPTIMIZER,
              metrics=['accuracy'])

model.summary()

model.fit(X, y, epochs=EPOCHS, batch_size=BATCH_SIZE)

model.save('out.model')
```



Encoding the **y** Vector

In the previous examples, we used `categorical_crossentropy`. This requires that all the targets are encoded with one-hot vectors. For instance:

- determiner: [1, 0, 0, 0]
- noun: [0, 1, 0, 0]
- verb: [0, 0, 1, 0]
- adjective: [0, 0, 0, 1]

With `sparse_categorical_crossentropy`, we can use numerical indices:

- determiner: 1
- noun: 2
- verb: 3
- adjective: 4

We do not need to use the `to_categorical` function.



Preprocessing

Preprocessing is more complex though: Four steps:

- 1 Read the corpus

```
train_sentences, dev_sentences, test_sentences, \
    column_names = load_ud_en_ewt()
```

- 2 Store the rows of the CoNLL corpus in dictionaries

```
conll_dict = CoNLLDictorizer(column_names, col_sep='\t')
train_dict = conll_dict.transform(train_sentences)
test_dict = conll_dict.transform(test_sentences)
```

- 3 Extract the features and store them in dictionaries

```
context_dictorizer = ContextDictorizer()
context_dictorizer.fit(train_dict)
X_dict, y_cat = context_dictorizer.transform(train_dict)
```

- 4 Vectorize the symbols

```
# We transform the X symbols into numbers
dict_vectorizer = DictVectorizer()
X_num = dict_vectorizer.fit_transform(X_dict)
```



Code Example

Jupyter Notebook: <https://github.com/pnugues/edan95/blob/master/programs/4.1-nn-pos-tagger.ipynb>




Architecture 2: Dense Word Encoding

One-hot encoding is results in sparse matrices

An alternative is to use dense vectors, for instance from a principal component analysis. Compare the character encodings of:

Sing, O Goddess, the anger...



Notebook, where we use GloVe:

ilppp/programs/ch08/python/lr_pos_tagger_emb.ipynb



Architecture 2: Using Embeddings

We replace the one-hot vectors with embeddings, the rest being the same. Word embeddings are dense vectors obtained by a principal component analysis or another method.

They can be trained by the neural network or pretrained

In this implementation:

- 1 We use pretrained embeddings from the GloVe project;
- 2 Our version of GloVe is lowercased, so we set all the characters in lowercase;
- 3 We add the embeddings as an `Embedding` layer at the start of the network;
- 4 We initialize the embedding layer with GloVe and make it trainable or not.

It would be possible to use a randomly initialized matrix as embeddings instead



The Embedding Layer

```
embedding_matrix = np.random.random(  
    (len(vocabulary_words) + 1,  
     EMBEDDING_DIM))  
  
for word in vocabulary_words:  
    if word in embeddings_dict:  
        embedding_matrix[word_idx[word]] = embeddings_dict[word]  
  
model = models.Sequential([  
    Embedding(cnt_uniq,  
              EMBEDDING_DIM,  
              trainable=True,  
              embeddings_initializer=  
                  initializers.Constant(embedding_matrix),  
              input_length=2 * W_SIZE + 1),  
    Flatten(),  
    Dense(NB_CLASSES, activation='softmax')  
])
```



Code Example

Jupyter Notebook: <https://github.com/pnugues/edan95/blob/master/programs/4.2-nn-pos-tagger-embeddings.ipynb>



Recap

Recap of the two architectures we saw in the previous lecture (September 15, 2022)




Recap: Feed-Forward Structure

- As input, the classifier uses:
 $(w_{i-2}, w_{i-1}, w_i, w_{i+1}, w_{i+2})$ to predict the part-of-speech tag t_i at index i .
- Here:
 (countrymen, to, visit, Mexico, ",")
 to predict VB

ID	FORM	PPOS	
	BOS	BOS	Padding
	BOS	BOS	
1	Battle	NN	
2	-	HYPH	
3	tested	NN	
...	
17	the	DT	
18	first	JJ	
19	of	IN	
20	their	PRP\$	
21	countrymen	NNS	Input features
22	to	TO	
23	visit	VB	Predicted tag
24	Mexico		↓
25	,		
26	a		
27	boatload		
...	
34	years		
35	ago		
36	.		
	EOS		Padding
	EOS		




Recap: Feed Forward (Multinomial)



Output: Softmax to predict the parts of speech



Recap: Feed Forward (Multinomial)



Input: Two kinds of **vectorization**: one-hot encoding and embeddings (GloVe in the examples)

Output: Softmax to predict the parts of speech



Architecture 3: Recurrent Structure

- In 2000, the CoNLL winners used *dynamic tags*
- The feature vector incorporates past predictions:
 $(w_{i-2}, w_{i-1}, w_i, w_{i+1}, w_{i+2}, t_{i-2}, t_{i-1})$,
 to predict the part-of-speech tag t_i at index i .
- Here:
 (countrymen, to, visit, Mexico, ",",
 NNS, TO)
 to predict VB

ID	FORM	PPOS	
	BOS	BOS	Padding
	BOS	BOS	
1	Battle	NN	
2	-	HYPH	
3	tested	NN	
...	
17	the	DT	
18	first	JJ	
19	of	IN	
20	their	PRP\$	
21	countrymen	NNS	Input features
22	to	TO	
23	visit	VB	Predicted tag
24	Mexico		↓
25	,		
26	a		
27	boatload		
...	
34	years		
35	ago		
36	.		
	EOS		Padding
	EOS		




Feature Vectors

ID	Feature vectors: X							PPOS: y
	w_{i-2}	w_{i-1}	w_i	w_{i+1}	w_{i+2}	t_{i-2}	t_{i-1}	
1	BOS	BOS	Battle	-	tested	BOS	BOS	NN
2	BOS	Battle	-	tested	Japanese	BOS	NN	HYPH
3	Battle	-	tested	Japanese	industrial	NN	HYPH	JJ
...
19	the	first	of	their	countrymen	DT	JJ	IN
20	first	of	their	countrymen	to	JJ	IN	PRP\$
21	of	their	countrymen	to	visit	IN	PRP\$	NNS
22	their	countrymen	to	visit	Mexico	PRP\$	NNS	TO
23	countrymen	to	visit	Mexico	,	NNS	TO	VB
24	to	visit	Mexico	,	a	TO	VB	NNP
25	visit	Mexico	,	a	boatload	VB	NNP	,
...
34	ashore	375	years	ago	.	RB	CD	NNS
35	375	years	ago	.	EOS	CD	NNS	RB
36	years	ago	.	EOS	EOS	NNS	RB	.




The RNN Architecture



A simple recurrent neural network; the dashed lines represent trainable connections.



The RNN Architecture (Unfolded)



The network unfolded in time. Equation used by implementations¹.

$$\mathbf{y}_{(t)} = \tanh(\mathbf{W}\mathbf{x}_{(t)} + \mathbf{U}\mathbf{y}_{(t-1)} + \mathbf{b})$$



¹See: <https://pytorch.org/docs/stable/nn.html#torch.nn.RNN>

Input Format for RNNs

The input format is different from feed forward networks.

We need to build two lists: one for the input and the other for the output

y	DET	NOUN	VERB	DET	NOUN
x	The	waiter	brought	the	meal

All the vectors in a same batch must have the same length. We pad them:

y	PAD	PAD	PAD	DET	NOUN	VERB	DET	NOUN
x	PAD	PAD	PAD	The	waiter	brought	the	meal

We could apply the padding after



Building the Sequences

```
def build_sequences(corpus_dict, key_x='form', key_y='pos',
                    tolower=True):
    X, Y = [], []
    for sentence in corpus_dict:
        x, y = [], []
        for word in sentence:
            x += [word[key_x]]
            y += [word[key_y]]
        if tolower:
            x = list(map(str.lower, x))
        X += [x]
        Y += [y]
    return X, Y
```

At this point, we have **x** and **y** vectors of symbols



Building Index Sequences

0 is for the padding symbol and 1 for the unknown words

```
idx_word = dict(enumerate(vocabulary_words, start=2))  
idx_pos = dict(enumerate(pos, start=1))  
word_idx = {v: k for k, v in idx_word.items()}  
pos_idx = {v: k for k, v in idx_pos.items()}
```

At this point, we have **x** and **y** vectors of numbers



Padding the Index Sequences

We build the complete **X_idx** and **Y_idx** matrices for the whole corpus
And we pad the matrices:

```
X = pad_sequences(X_idx)
```

```
Y = pad_sequences(Y_idx)
```

```
# The number of POS classes and 0 (padding symbol)
```

```
Y_train = to_categorical(Y, num_classes=len(pos) + 1)
```

`pad_sequences` can have an argument that specifies the maximal length
`maxlen (MAX_SEQUENCE_LENGTH)`.

The padded sentences must have the same length in a batch. This is
automatically computed by Keras



Recurrent Neural Networks (RNN)

```
model = models.Sequential([  
    Embedding(len(vocabulary_words) + 2,  
              EMBEDDING_DIM,  
              mask_zero=True,  
              embeddings_initializer=  
                  initializers.Constant(embedding_matrix),  
              trainable=True,  
              input_length=None),  
    SimpleRNN(100, return_sequences=True),  
    # Bidirectional(SimpleRNN(100, return_sequences=True)),  
    Dense(NB_CLASSES + 1, activation='softmax')])
```



Parameters

Keras functions have many parameters.

In case of doubt, read the documentation

A few useful parameters:

- 1 `mask_zero=True` is to tell whether or not the input value 0 is a special “padding” value;
- 2 `return_sequences=True` tells whether to return the last output in the output sequence, or the full sequence. In sequences, it is essential;
- 3 `recurrent_dropout=0.3` tells how much to drop for the linear transformation of the recurrent state.



Code Example

Jupyter Notebook: <https://github.com/pnugues/edan95/blob/master/programs/4.3-rnn-pos-tagger.ipynb>



LSTMs

Simple RNNs use the previous output as input. They have then a very limited feature context.

Long short-term memory units (LSTM) are an extension to RNNs that can remember, possibly forget, information from longer or more distant sequences.

Given an input at index t , \mathbf{x}_t , a LSTM unit produces:

- A short term state, called \mathbf{h}_t and
- A long-term state, called \mathbf{c}_t or memory cell.

The short-term state, \mathbf{h}_t , is the unit output, i.e. \mathbf{y}_t ; but both the long-term and short-term states are reused as inputs to the next unit.



LSTM Equations

A LSTM unit starts from a core equation that is identical to that of a RNN:

$$\mathbf{g}_t = \tanh(\mathbf{W}_g \mathbf{x}_t + \mathbf{U}_g \mathbf{h}_{t-1} + \mathbf{b}_g).$$

From the previous output and current input, we compute three kinds of filters, or gates, that will control how much information is passed through the LSTM cell

The two first gates, \mathbf{i} and \mathbf{f} , defined as:

$$\begin{aligned}\mathbf{i}_t &= \text{activation}(\mathbf{W}_i \mathbf{x}_t + \mathbf{U}_i \mathbf{h}_{t-1} + \mathbf{b}_i), \\ \mathbf{f}_t &= \text{activation}(\mathbf{W}_f \mathbf{x}_t + \mathbf{U}_f \mathbf{h}_{t-1} + \mathbf{b}_f),\end{aligned}$$

model respectively how much we will keep from the base equation and how much we will forget from the long-term state.



LSTM Equations (II)

To implement this selective memory, we apply the two gates to the base equation and to the previous long-term state with the element-wise product (Hadamard product), denoted \circ , and we sum the resulting terms to get the current long-term state:

$$\mathbf{c}_t = \mathbf{i}_t \circ \mathbf{g}_t + \mathbf{f}_t \circ \mathbf{c}_{t-1}.$$

The third gate:

$$\mathbf{o}_t = \text{activation}(\mathbf{W}_o \mathbf{x}_t + \mathbf{U}_o \mathbf{h}_{t-1} + \mathbf{b}_o)$$

modulates the current long-term state to produce the output:

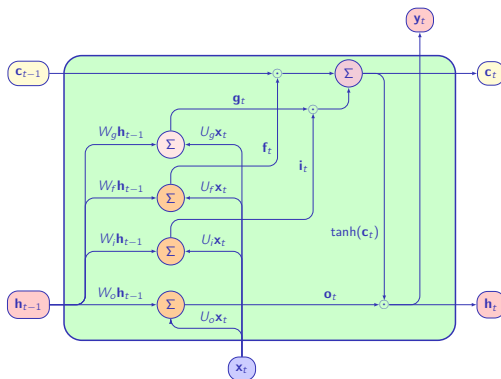
$$\mathbf{h}_t = \mathbf{o}_t \circ \tanh(\mathbf{c}_t).$$

The LSTM parameters are determined by a gradient descent.
See also:

<https://pytorch.org/docs/stable/nn.html#torch.nn.LSTM>



The LSTM Architecture



An LSTM unit showing the data flow, where \mathbf{g}_t is the unit input, \mathbf{i}_t , the input gate, \mathbf{f}_t , the forget gate, and \mathbf{o}_t , the output gate. The activation functions have been omitted



Long Short-Term Memory (LSTM)

```
model = models.Sequential([
    Embedding(len(vocabulary_words) + 2,
              EMBEDDING_DIM,
              mask_zero=True,
              embeddings_initializer=
                  initializers.Constant(embedding_matrix),
              trainable=True,
              input_length=None),
    Bidirectional(LSTM(100, return_sequences=True)),
    Dense(NB_CLASSES + 1, activation='softmax')])
```



Course Content: 2022

The rest of the slides in this document is not part of the course in 2022.

You can nonetheless read it if you are curious.



Conditional Random Fields

A tagger produces a sequence of tags, where a given tag obviously depends on the previous ones.

For instance, a preposition cannot follow a determiner

We can model the tag transitions probabilities using **conditional random fields** (CRF)

The simplest form is the **linear chain**.

If \mathbf{y} denotes the output, here a sequence of tags, and \mathbf{x} , a sequence of inputs, consisting for instance of the words and the characters, we try to maximize

$$P(\mathbf{y}|\mathbf{x}),$$

i.e.

$$\hat{\mathbf{y}} = \arg \max_{\mathbf{y}} P(\mathbf{y}|\mathbf{x}).$$



Conditional Random Fields (II)

As we want the output sequence to depend on the input and on previously predicted output, we rewrite $P(\mathbf{y}|\mathbf{x})$ as a joint probability, $P(\mathbf{y}, \mathbf{x})$, and, more precisely, as:


$$\begin{aligned}
 P(\mathbf{y}|\mathbf{x}) &= \frac{P(\mathbf{y}, \mathbf{x})}{\sum_{\mathbf{y}' \in Y} P(\mathbf{y}', \mathbf{x})}, \\
 &= \frac{\prod \exp(w_1 \cdot f(y_t, \mathbf{x}, t)) \cdot \prod \exp(w_2 \cdot f(y_t, y_{t+1}))}{\sum_{\mathbf{y}' \in Y} P(\mathbf{y}', \mathbf{x})},
 \end{aligned}$$

where \mathbf{y}' denotes a sequence, Y , the set of all the possible sequences, and $\sum_{\mathbf{y}' \in Y} P(\mathbf{y}', \mathbf{x})$ is a normalizing factor to have a sum of probabilities of one



Viterbi (Informal)

Je le donne demain dans la matinée 'I give it tomorrow in the morning'



Viterbi (Informal)

The term brought by the word *demain* has still the memory of the ambiguity of *donne*: $P(\text{adv}|\text{verb}) \times P(\text{demain}|\text{adv})$ and $P(\text{adv}|\text{noun}) \times P(\text{demain}|\text{adv})$.

This is no longer the case with *dans*.

According to the noisy channel model and the bigram assumption, the term brought by the word *dans* is $P(\text{dans}|\text{prep}) \times P(\text{prep}|\text{adv})$.

It does not show the ambiguity of *le* and *donne*.


The subsequent terms will ignore it as well.

We can discard the corresponding paths.

The optimal path does not contain nonoptimal subpaths.



Trellis Representation



Filling the Trellis

$i \backslash \delta$	δ_1	δ_2	δ_3	δ_4	δ_5	δ_6	δ_7	δ_8
PREP	0							
ADV	0							
PRO	0							
VERB	0							
NOUN	0							
ART	0							
<s>	1.0	0	0	0	0	0	0	0
	<s>	<i>Je</i>	<i>le</i>	<i>donne</i>	<i>demain</i>	<i>dans</i>	<i>la</i>	<i>matinée</i>

To fill the δ_3 column, for each cell j , we compute

$$\max_i P(j|i) \times P(le|j) \times \delta_2(i).$$

The pronoun cell, for instance, is filled with

$$\max_i P(\text{PRO}|i) \times P(le|\text{PRO}) \times \delta_2(i).$$



Supervised Learning: A Summary

Needs a manually annotated corpus called the **Gold Standard**

The Gold Standard may contain errors (*errare humanum est*) that we ignore

A classifier is trained on a part of the corpus, the **training set**, and evaluated on another part, the **test set**, where automatic annotation is compared with the “Gold Standard”

N-fold cross validation is used avoid the influence of a particular division
Some algorithms may require additional optimization on a development set

Classifiers can use statistical or symbolic methods

